

Roteiro da implantação e teste

As características do computador usado para realizar a implantação e teste do procedimento detalhado neste documento foram:

- Modelo do CPU: Apple M1
- Número de cores: 8
- Memória RAM: 16 Gb

1. Pré-requisitos

Clone o [repositório do GitHub](#) com os arquivos que serão usados para a implantação.

Instalar os seguintes pacotes:

- minikube
- kubectl
- aws-cli

```
→ ~ minikube version
minikube version: v1.33.1
commit: 5883c09216182566a63dff4c326a6fc9ed2982ff
→ ~ kubectl version --client
Client Version: v1.30.3
Kustomize Version: v5.0.4-0.20230601165947-6ce0bf390ce3
→ ~ aws --version
aws-cli/2.17.17 Python/3.11.9 Darwin/23.5.0 source/arm64
```

2. Cluster Kubernetes usando minikube

Iniciamos o Minikube

```
$ minikube start --cpus 2 --memory 4096 --driver=docker
```

```
→ ~ minikube start --cpus 2 --memory 4096 --driver=docker
🐳 minikube v1.33.1 en Darwin 14.5 (arm64)
🌟 Using the docker driver based on existing profile
👍 Starting "minikube" primary control-plane node in "minikube" cluster
📦 Pulling base image v0.0.44 ...
🔄 Restarting existing docker container for "minikube" ...
🔧 Preparando Kubernetes v1.30.0 en Docker 26.1.1...
🔍 Verifying Kubernetes components...
  ■ Using image gcr.io/k8s-minikube/storage-provisioner:v5
  ■ Using image registry.k8s.io/metrics-server/metrics-server:v0.7.1
  ■ Using image docker.io/kubernetes/dashboard:v2.7.0
  ■ Using image docker.io/kubernetes/metrics-scraper:v1.0.8
💡 Some dashboard features require the metrics-server addon. To enable all features please run:

    minikube addons enable metrics-server

🌟 Complementos habilitados: metrics-server, storage-provisioner, default-storageclass, dashboard
🏠 Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
→ ~ minikube addons enable metrics-server
```

Habilitamos as métricas

```
$ minikube addons enable metrics-server
```

```
→ ~ minikube addons enable metrics-server
💡 metrics-server is an addon maintained by Kubernetes. For any concerns contact minikube on GitHub.
You can view the list of minikube maintainers at: https://github.com/kubernetes/minikube/blob/master/OWNERS
▪ Using image registry.k8s.io/metrics-server/metrics-server:v0.7.1
🌟 The 'metrics-server' addon is enabled
```

Criamos um deployment usando o arquivo [deployment.yaml](#)

```
$ kubectl apply -f deployment.yaml
```

Exporemos o serviço usando o arquivo [service.yaml](#)

```
$ kubectl apply -f service.yaml
```

Configurar o autoescalamento horizontal usando o arquivo [hpa.yaml](#)

```
$ kubectl apply -f hpa.yaml
```

```
→ src kubectl apply -f deployment.yaml
deployment.apps/web-server unchanged
→ src kubectl apply -f service.yaml
service/web-server-service unchanged
→ src kubectl apply -f hpa.yaml
horizontalpodautoscaler.autoscaling/web-server-hpa unchanged
```

Verificamos o deployment

```
$ kubectl get pods
$ kubectl get svc
$ kubectl get hpa
```

```
→ ~ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
web-server-d7f8d6c6-8hcv7           1/1     Running   0           117s
web-server-d7f8d6c6-jf9h1           1/1     Running   0           117s
web-server-d7f8d6c6-jmkm2           1/1     Running   0           117s
→ ~ kubectl get hpa
NAME      REFERENCE          TARGETS          MINPODS   MAXPODS   REPLICAS   AGE
web-server-hpa  Deployment/web-server  cpu: <unknown>/50%    1         5         3         2m3s
→ ~ kubectl get svc
NAME                TYPE          CLUSTER-IP      EXTERNAL-IP   PORT(S)          AGE
kubernetes          ClusterIP     10.96.0.1       <none>        443/TCP          17d
web-server-service  LoadBalancer 10.103.87.138   <pending>     80:32145/TCP     2m36s
→ ~
```

Geramos carga no cluster para testar o HPA

```
$ minikube ssh
docker@minikube:~$ sudo apt-get update
docker@minikube:~$ sudo apt-get install stress -y
```

```

docker@minikube:~$ sudo apt-get install stress -y
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following NEW packages will be installed:
  stress
0 upgraded, 1 newly installed, 0 to remove and 38 not upgraded.
Need to get 18.2 kB of archives.
After this operation, 47.1 kB of additional disk space will be used.
Get:1 http://ports.ubuntu.com/ubuntu-ports jammy/universe arm64 stress arm64 1.0.5-1 [18.2 kB]
Fetched 18.2 kB in 2s (11.8 kB/s)
debconf: delaying package configuration, since apt-utils is not installed
Selecting previously unselected package stress.
(Reading database ... 10937 files and directories currently installed.)
Preparing to unpack .../stress_1.0.5-1_arm64.deb ...
Unpacking stress (1.0.5-1) ...
Setting up stress (1.0.5-1) ...

```

Geramos carga no cluster para testar o HPA

```

$ kubectl run -i --tty load-generator --image=busybox /bin/sh
# while sleep 0.01; do wget -q -O- http://web-server-service; done

```

```

→ ~ kubectl run -i --tty load-generator --image=busybox /bin/sh
If you don't see a command prompt, try pressing enter.
/ #
/ #
/ # █

```

Verificamos o autoescalamento horizontal no dashboard

```

$ minikube dashboard

```

```

→ ~ minikube dashboard
🤖 Verifying dashboard health ...
🚀 Launching proxy ...
🤖 Verifying proxy health ...
🌐 Opening http://127.0.0.1:53335/api/v1/namespaces/kubernetes-dashboard/services/http:kubernetes-dashboard:/proxy/ in your default browser...
^C

```



3. Cluster AWS EKS

Criar o perfil IAM com as políticas de gerenciamento. Usar o arquivo [eks-cluster-role-trust-policy.json](#) e executar os seguintes comandos

```
$ aws iam create-role \  
  --role-name myAmazonEKSClusterRole \  
  --assume-role-policy-document file://"eks-cluster-role-trust-policy.json"  
  
$ aws iam attach-role-policy \  
  --policy-arn arn:aws:iam::aws:policy/AmazonEKSClusterPolicy \  
  --role-name myAmazonEKSClusterRole
```

```

→ src nano eks-cluster-role-trust-policy.json
→ src
{
  "Role": {
    "Path": "/",
    "RoleName": "myAmazonEKSClusterRole",
    "RoleId": "AROAZUC3F462LGKLPVCPU",
    "Arn": "arn:aws:iam::730335668148:role/myAmazonEKSClusterRole",
    "CreateDate": "2024-08-12T23:57:08+00:00",
    "AssumeRolePolicyDocument": {
      "Version": "2012-10-17",
      "Statement": [
        {
          "Effect": "Allow",
          "Principal": {
            "Service": "eks.amazonaws.com"
          },
          "Action": "sts:AssumeRole"
        }
      ]
    }
  }
}
...skipping...
→ src aws iam create-role \
--role-name myAmazonEKSClusterRole \
--assume-role-policy-document file://"eks-cluster-role-trust-policy.json"
→ src aws iam attach-role-policy \
--policy-arn arn:aws:iam::aws:policy/AmazonEKSClusterPolicy \
--role-name myAmazonEKSClusterRole

```

Criamos o cluster e nós no AWS EKS usando o seguinte comando

```
$ aws eks update-kubeconfig --region us-east-1 --name my-cluster
```

```

2024-07-29 17:26:17 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-ng-4557568c"
2024-07-29 17:26:47 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-ng-4557568c"
2024-07-29 17:27:25 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-ng-4557568c"
2024-07-29 17:28:13 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-ng-4557568c"
2024-07-29 17:29:46 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-ng-4557568c"
2024-07-29 17:29:46 [i] waiting for the control plane to become ready
2024-07-29 17:29:47 [✓] saved kubeconfig as "/Users/fabiotorres/.kube/config"
2024-07-29 17:29:47 [i] no tasks
2024-07-29 17:29:47 [✓] all EKS cluster resources for "my-cluster" have been created
2024-07-29 17:29:47 [✓] created 0 nodegroup(s) in cluster "my-cluster"
2024-07-29 17:29:48 [i] nodegroup "ng-4557568c" has 2 node(s)
2024-07-29 17:29:48 [i] node "ip-192-168-3-224.ec2.internal" is ready
2024-07-29 17:29:48 [i] node "ip-192-168-44-4.ec2.internal" is ready
2024-07-29 17:29:48 [i] waiting for at least 1 node(s) to become ready in "ng-4557568c"
2024-07-29 17:29:48 [i] nodegroup "ng-4557568c" has 2 node(s)
2024-07-29 17:29:48 [i] node "ip-192-168-3-224.ec2.internal" is ready
2024-07-29 17:29:48 [i] node "ip-192-168-44-4.ec2.internal" is ready
2024-07-29 17:29:48 [✓] created 1 managed nodegroup(s) in cluster "my-cluster"
2024-07-29 17:29:49 [i] kubectl command should work with "/Users/fabiotorres/.kube/config", try 'kubectl get nodes'
2024-07-29 17:29:49 [✓] EKS cluster "my-cluster" in "us-east-1" region is ready

```

Verificar os pods do EKS criado

```
$ kubectl get pods -A -o wide
```

```
→ src git:(main) * kubectl get pods -A -o wide
NAMESPACE   NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE                                NOMINATED NODE   READINESS GATES
kube-system  aws-node-tmj6g                      2/2     Running   0           7m25s  192.168.43.199  ip-192-168-43-199.ec2.internal     <none>           <none>
kube-system  aws-node-vwfn2                      2/2     Running   0           7m35s  192.168.7.167   ip-192-168-7-167.ec2.internal      <none>           <none>
kube-system  coredns-586b798467-7m5v8           1/1     Running   0           11m    192.168.27.216  ip-192-168-7-167.ec2.internal      <none>           <none>
kube-system  coredns-586b798467-mpp5j           1/1     Running   0           11m    192.168.28.231  ip-192-168-7-167.ec2.internal      <none>           <none>
kube-system  kube-proxy-gl7v7                    1/1     Running   0           7m25s  192.168.43.199  ip-192-168-43-199.ec2.internal     <none>           <none>
kube-system  kube-proxy-jhvpj                    1/1     Running   0           7m35s  192.168.7.167   ip-192-168-7-167.ec2.internal      <none>           <none>
```

Visualizar o uso de recursos com o Kubernetes Metrics Server

```
$ kubectl apply -f
https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
$ kubectl get deployment metrics-server -n kube-system
```

```
→ src git:(main) * kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
serviceaccount/metrics-server created
clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created
clusterrole.rbac.authorization.k8s.io/system:metrics-server created
rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader created
clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator created
clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server created
service/metrics-server created
deployment.apps/metrics-server created
apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io created
→ src git:(main) * kubectl get deployment metrics-server -n kube-system
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
metrics-server      1/1     1             1           53s
```

Configurar o Horizontal Pod Autoscaler

```
$ kubectl apply -f https://k8s.io/examples/application/php-apache.yaml
$ kubectl autoscale deployment php-apache --cpu-percent=50 --min=1 --max=10
$ kubectl get hpa
```

```
→ src git:(main) * kubectl apply -f https://k8s.io/examples/application/php-apache.yaml
deployment.apps/php-apache created
service/php-apache created
→ src git:(main) * kubectl autoscale deployment php-apache --cpu-percent=50 --min=1 --max=10
horizontalpodautoscaler.autoscaling/php-apache autoscaled
```

```
→ src git:(main) * kubectl get hpa
NAME           REFERENCE                TARGETS      MINPODS   MAXPODS   REPLICAS   AGE
php-apache     Deployment/php-apache     cpu: 0%/50%  1         10        1           42s
```

Teste do autoescalamento

```
$ kubectl run -i \
  --tty load-generator \
  --rm --image=busybox \
  --restart=Never \
  -- /bin/sh -c "while sleep 0.01; do wget -q -O- http://php-apache; done"
```

[illegible]

```
→ ~ kubectl get hpa
NAME                REFERENCE                TARGETS          MINPODS   MAXPODS   REPLICAS   AGE
php-apache          Deployment/php-apache     cpu: 150%/50%    1          10         1          104s
→ ~ kubectl get hpa php-apache
NAME                REFERENCE                TARGETS          MINPODS   MAXPODS   REPLICAS   AGE
php-apache          Deployment/php-apache     cpu: 251%/50%    1          10         3          113s
→ ~ kubectl get hpa
NAME                REFERENCE                TARGETS          MINPODS   MAXPODS   REPLICAS   AGE
php-apache          Deployment/php-apache     cpu: 72%/50%     1          10         9          2m36s
→ ~
```